

✔ Congratulations! You passed!

Grade received 100% To pass 80% or higher

Go to next item

Module Review

Latest Submission Grade 100%

1. Complete the following:

1 / 1 point

You should feed your machine learning model your _____ and not your _____. It will learn those for itself!

- ☒ data, rules
- ☐ If/then statements, data
- ☐ rules, data

✔ Correct
Correct!

2. True or False: Cloud SQL is a big data analytics warehouse

1 / 1 point

- ☐ True
- ☒ False

✔ Correct
Correct - Cloud SQL is a transaction RDBMS or relational database management system. It is designed for many more WRITES than READS.

Whereas BigQuery is a big data analytics warehouse which is optimized for reporting READS.

3. True or False: If you are migrating your Hadoop workload to the cloud, you must first rewrite all your Spark jobs to be compliant with the cloud.

1 / 1 point

- ☐ True
- ☒ False

✔ Correct
Correct - you can run your same Spark job code running on the same Hadoop software but running on cloud hardware with Cloud Dataproc.

4. You are thinking about migrating your Hadoop workloads to the cloud and you have a few workloads that are fault-tolerant (they can handle interruptions of individual VMs gracefully). What are some architecture considerations you should explore in the cloud? Choose all that apply

1 / 1 point

☒ Use PVMs or Preemptible Virtual Machines

✔ Correct
Correct!

☒ Migrate your storage from on-cluster HDFS to off-cluster Google Cloud Storage (GCS)

✔ Correct
Correct!

☒ Consider having multiple Cloud Dataproc instances for each priority workload and then turning them down when not in use

☒ **Correct**
Correct!

5. Google Cloud Storage is a good option for storing data that:

1 / 1 point

(Select the 2 correct options below).

☐ Is ingested in real-time from sensors and other devices and supports SQL-based queries

☒ May be required to be read at some later time (i.e. load a CSV file into BigQuery)

☒ **Correct**
Correct!

☐ Will be accessed frequently and updated constantly with new transactions from a front-end and needs to be stored in a relational database

☒ May be imported from a bucket into a Hadoop cluster for analysis

☒ **Correct**
Correct!

6. Relational databases are a good choice when you need:

1 / 1 point

☐ Aggregations on unstructured data

☐ Fast queries on terabytes of data

☐ Streaming, high-throughput writes

☒ Transactional updates on relatively small datasets

☒ **Correct**
Correct!

7. Cloud SQL and Cloud Dataproc offer familiar tools (MySQL and Hadoop/Pig/Hive/Spark). What is the value-add provided by Google Cloud Platform?

1 / 1 point

(Select the 2 correct options below)

☒ Running it on Google infrastructure offers reliability and cost savings

☒ **Correct**
Yes. You pay only for the resources you use. Cloud SQL can be shut down when it's not being used. Hadoop clusters can be of preemptible nodes, and so on.

☐ It's the same API, but Google implements it better

☒ Fully-managed versions of the software offer no-ops

☒ **Correct**
Yes. No-ops is the main value-add here.

☐ Google-proprietary extensions and bug fixes to MySQL, Hadoop, and so on